

This is an experiment on using probabilistic programming (using Soss.jl and Turing.jl) to calibrate review scores given by n_r reviewers to n_a articles submitted to, e.g., a conference for publication. Each reviewer assigns each article a score between 1 and 11 (categories start from 1, otherwise it would have been 0-10). We assume there are no missing reviews for simplicity. The review scores are stored in the vector R_v . The reviewers might not use a "metric scale" for their scores, i.e., difference between score 3-4 might be larger or smaller than the difference between 9-10. Due to this, we make use of ordinal regression, where we estimate the scale the reviewers are using. In reality, we might want to have a separate scale for each reviewer, but the number of reviews per reviewer required to estimate this accurately would be too high, so we settle for one single scale for all of them.

```
cd(@_DIR_)
using Pkg
pkg"activate ."

Activating environment at `~/soss_reviews/Project.toml`

using Soss, MonteCarloMeasurements, Distributions, Turing, Plots, LinearAlgebra, Statistics
default(size=(500,300))
nr = 10 # Number of reviewers
na = 10 # Number of articles
indv = [(i,j) for i in 1:nr, j in 1:na]
nscores = 11;
```

Below, we define some helper functions

```
function Distributions.cdf(R::AbstractArray,mi=minimum(skipmissing(R)),ma=maximum(skipmissing(R)))
    range = mi:ma
    Ro = copy(R)
    cumdist = map(range) do level
        count(R .<= level) / length(R)
    end
    cumdist
end

function rankdist(x,y)
    d = 0
    for xi in eachindex(x)
        yi = findfirst(==(x[xi]), y)
        d += abs(xi-yi)
    end
    d/length(x)
end

function logodds(R::AbstractArray)
    cumdist = cdf(R)
    mi,ma = extrema(R)
    range = mi:ma
    Ro = copy(R)
    map(R) do r
        ind = findfirst(==(r), range)
        logit(cumdist[ind])
    end
end;
```

This function accepts a vector of differences between cutpoints and forms the cumulative sum in both directions to form the final cutpoints. The reason for calculating the cumulative sum in both directions is that the variance grows as you add uncertain variables, and with this approach, the variance will be highest in the middle, which agrees with my intuition on how flexible the cutpoints should be.

```
cumcut(diffcp) = ((cumsum(diffcp) + reverse(1 .- cumsum(reverse(diffcp)))) ./ 2)*12 .- 6;
```

1 Soss

First out is the Soss model

```
cum_model = Soss.@model indv begin
    rσ ~ Gamma(0.2) # The variance of the noise each reviewer has in their review is drawn from a common pool
    article_pop_std ~ truncated(Normal(1., 0.1), 0, 100) # The population of all article scores has a common variance
    reviewer_noise ~ truncated(Normal(rσ, 0.1), 0, 3) |> iid(nr) # Different reviewer have different noise variances
    reviewer_gain ~ Normal(1, 0.15) |> iid(nr) # Each reviewer has a unique gain, i.e., when an article gets better or worse, the reviewer adjusts the score more or less.
    article_score ~ Normal(0,article_pop_std) |> iid(na) # The true, calibrated article score
    diffcp ~ Dirichlet(nscores-1,50) # Vector of differences between cutpoints
```

```

cutpoints = cumcut(diffcp)

z ~ Normal(0,1) > iid(length(indv)) # Vector of noise components, one for each review
Rv ~ For(length(indv)) do ind
  i,j = indv[ind]
  pred = article_score[j] + reviewer_noise[i]*z[ind] + reviewer_gain[i]*article_score[j] # linear model predicting the log-odds of the review score, this will be roughly
  # between -4 and 4
  OrderedLogistic(pred,cutpoints) # The observed review score is an ordered logistic variable. This transform the `pred` to a categorical value between 1 and 10
end
end;

```

below, we sample one data point from the model and call this the "truth".

```

s = [rand(cum_model(indv=indv)) for _ in 1:1000] # rand(::SossModel) does not seem to accept a number of samples
truth = rand(cum_model(indv=indv));

```

We now perform inference on the model using as observed variables the review scores from the "truth"

```
@time post = dynamicHMC(cum_model(indv=indv), (Rv=truth.Rv.), 1000);
```

194.804423 seconds (1.08 G allocations: 147.633 GiB, 17.75% gc time)

This call takes a long time. It seems to be taking long time before it even starts to sample.

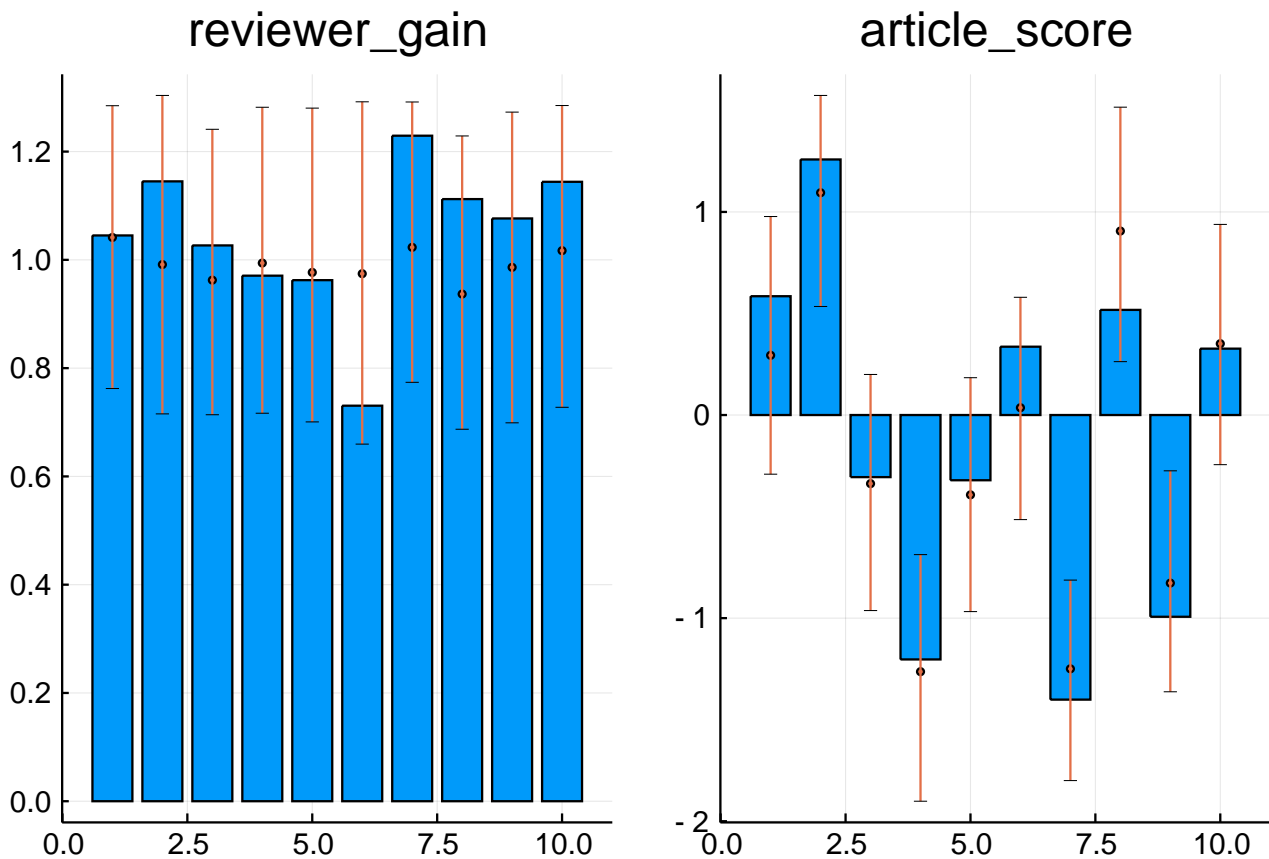
```
p = particles(post[200:end]); # Convert posterior to `Particles`
```

We now visualize the inferred articles scores and reviewer gains and compare the posterior to what we know are the generating parameters from the truth

```

figs = map( (:reviewer_gain, :article_score) ) do s
  bar(getproperty(truth, s))
  prop = getproperty(p, s)
  errorbarplot!(1:length(prop), prop, seriestype=:scatter, legend=false, title=string(s), m=2)
end
plot(figs...)

```



using Soss, there is a clear indication that the model has learned at least something about the reviewer gains, even though the variance is high.

We can calculate the log-odds of the observed scores for each article and compare this to the models predictions

```

lo = logodds(truth.Rv)
observed_score = map(1:na) do j
  median([lo[ind] for ind in eachindex(indv) if indv[ind][2] == j])
end

function print_rank_results(truth, p, observed_score)


  println("Percentage of correct rank from model $(mean(sortperm(truth.article_score) .== sortperm(mean.(p.article_score))))")
  println("Percentage of correct rank from observed score $(mean(sortperm(truth.article_score) .== sortperm(observed_score))))")

  println("Rankdist between correct rank and model $(rankdist(sortperm(truth.article_score), sortperm(mean.(p.article_score))))") # `rankdist` is a distance measure I cooked
  up to try to measure the permutation distance between two rank vectors.
  println("Rankdist between correct rank and observed score $(rankdist(sortperm(truth.article_score), sortperm(observed_score))))")
  scatter([truth.article_score (observed_score)], label=["true" "obs"], title="Article scores", xlabel="Article number")
  errorbarplot!(1:na, p.article_score, seriestype=:scatter, lab="model", m=(2,))
end
print_rank_results(truth, p, observed_score)

```

Percentage of correct rank from model 0.4
 Percentage of correct rank from observed score 0.2
 Rankdist between correct rank and model 0.8
 Rankdist between correct rank and observed score 1.0

Article scores



2 Turing

We now perform the same exercise with Turing, the model should be exactly the same with the same numerical values for the parameters

```

using Turing2MonteCarloMeasurements, NamedTupleTools
Turing.@model cum_model(indv, Rv, ::Type{T}=Float64) where {T} = begin
  rσ ~ Gamma(0.2)
  article_pop_std ~ truncated(Normal(1., 0.1), 0, 100)
  reviewer_noise = Vector{T}(undef, nr)
  pred = Vector{T}(undef, length(indv))

  for i = 1:nr
    reviewer_noise[i] ~ truncated(Normal(rσ, 0.1), 0, 3)
  end
end

```

```

end
reviewer_gain ~ MvNormal(fill(1,nr), 0.15^2)
article_score ~ MvNormal(zeros(na),article_pop_std^2)

diffcp ~ Dirichlet(nscores-1,50)
cutpoints = cumcut(diffcp)
z ~ MvNormal(zeros(length(indv)),1)
for ind in eachindex(indv)
    i,j = indv[ind]
    pred[ind] = article_score[j] + reviewer_noise[i]*z[ind] + reviewer_gain[i]*article_score[j]
    Rv[ind] ~ OrderedLogistic(pred[ind],cutpoints)
end
@namedtuple(Rv, article_score, cutpoints, reviewer_noise, reviewer_gain, pred, diffcp, z, rσ, article_pop_std)
end;


```

Once again we sample one data points and call it the truth

```

prior = cum_model(indv, Union{Int,Missing}[fill(missing, length(indv))...])
# truth = prior() # We use the truth from Soss
prior_sample = [prior() for _ in 1:500] |> particles
errorbarplot(1:length(indv), prior_sample.Rv, 0.0, title="Review score prior") |> display

```




We can visualize what the prior considers possible cutpoints for the log-odds

```

mcplot(1:length(prior_sample.cutpoints), prior_sample.cutpoints, title="Cutpoints prior") |> display

```


Cutpoints prior



We can also visualize the prior distribution over observed review scores and log-odds


```
histogram(reduce(union,prior_sample.Rv), title="Samples of review scores from prior", xlabel="Review score")
```

Samples of review scores from prior



```
histogram(reduce(union,prior_sample.pred), title="Samples of log-odds predictions from prior")
```

Samples of log- odds predictions from prior



We now sample from the posterior using Turing

```
m = cum_model(indv, Int.(truth.Rv))
@time chain = sample(m, HMC(0.03, 7), 1200) # NUTS does not work
```

38.257058 seconds (185.22 M allocations: 23.894 GiB, 14.02% gc time)

```
p = Particles(chain, crop=200);
```

Turing samples *much* faster than Soss for this model

```
dc = describe(chain, digits=3, q=[0.1, 0.5, 0.9])
```

2-element Array{ChainDataFrame,1}

Summary Statistics

. Omitted printing of 1 columns|

Row	parameters	mean	std	naive_se	mcse	ess
1	Symbol	Float64	Float64	Float64	Float64	Any

1	article_pop_std	0.979	0.086	0.002	0.005	513.0
29						
2	article_score[1]	0.3	0.294	0.008	0.041	62.70
7						
3	article_score[2]	1.214	0.335	0.01	0.047	24.24
5						
4	article_score[3]	-0.468	0.297	0.009	0.048	19.40
8						
5	article_score[4]	-1.255	0.294	0.008	0.035	57.36
7						
6	article_score[5]	-0.29	0.273	0.008	0.032	74.41
9						
7	article_score[6]	0.139	0.327	0.009	0.053	10.23

8	article_score[7]	-1.301	0.331	0.01	0.036	47.0
9	article_score[8]	0.945	0.305	0.009	0.043	41.41
4						
10	article_score[9]	-0.828	0.303	0.009	0.044	28.17
8						
11	article_score[10]	0.42	0.285	0.008	0.033	73.23
2						
12	diffcp[1]	0.101	0.013	0.0	0.001	278.3
27						
13	diffcp[2]	0.101	0.013	0.0	0.001	266.5
22						
14	diffcp[3]	0.1	0.012	0.0	0.001	228.6
03						
15	diffcp[4]	0.106	0.013	0.0	0.001	354.1
2						
16	diffcp[5]	0.098	0.013	0.0	0.001	338.1
82						
17	diffcp[6]	0.095	0.012	0.0	0.001	247.6
84						
18	diffcp[7]	0.098	0.013	0.0	0.001	254.9
44						
19	diffcp[8]	0.101	0.013	0.0	0.001	237.5
54						
20	diffcp[9]	0.101	0.013	0.0	0.001	188.4
83						
21	diffcp[10]	0.097	0.013	0.0	0.002	88.62
9						
22	reviewer_gain[1]	1.002	0.021	0.001	0.001	925.1
23						
23	reviewer_gain[2]	0.999	0.023	0.001	0.0	874.9
01						
24	reviewer_gain[3]	0.999	0.023	0.001	0.001	905.6
79						
25	reviewer_gain[4]	1.0	0.021	0.001	0.001	930.4
1						
26	reviewer_gain[5]	1.0	0.022	0.001	0.0	1000.
94						
27	reviewer_gain[6]	0.999	0.023	0.001	0.0	1070.
27						
28	reviewer_gain[7]	1.001	0.022	0.001	0.001	1112.
77						
29	reviewer_gain[8]	0.998	0.023	0.001	0.001	437.7
74						
30	reviewer_gain[9]	1.0	0.022	0.001	0.001	859.7
32						
31	reviewer_gain[10]	0.999	0.023	0.001	0.001	884.9
1						
32	reviewer_noise[1]	0.787	0.222	0.006	0.054	10.11
5						
33	reviewer_noise[2]	0.8	0.236	0.007	0.059	9.589
34	reviewer_noise[3]	0.793	0.225	0.006	0.055	11.54
4						
35	reviewer_noise[4]	0.778	0.241	0.007	0.06	9.393
36	reviewer_noise[5]	0.775	0.227	0.007	0.055	9.89
37	reviewer_noise[6]	0.788	0.223	0.006	0.052	10.04
7						
38	reviewer_noise[7]	0.804	0.231	0.007	0.057	8.953
39	reviewer_noise[8]	0.79	0.226	0.007	0.055	10.00
5						
40	reviewer_noise[9]	0.804	0.227	0.007	0.056	10.63
5						
41	reviewer_noise[10]	0.788	0.226	0.007	0.056	9.435
42	rσ	0.786	0.211	0.006	0.056	9.637
43	z[1]	-0.241	0.751	0.022	0.164	17.80
9						
44	z[2]	0.585	1.008	0.029	0.229	13.71
5						
45	z[3]	-0.044	0.788	0.023	0.177	7.987

46	z[4]	-0.031	1.132	0.033	0.302	5.649
47	z[5]	-0.715	0.872	0.025	0.204	14.08
1	z[6]	-0.402	0.958	0.028	0.252	5.462
49	z[7]	0.296	0.546	0.016	0.105	24.99
50	z[8]	0.548	0.784	0.023	0.194	10.05
51	z[9]	0.54	0.772	0.022	0.165	24.93
52	z[10]	0.488	0.857	0.025	0.218	7.193
53	z[11]	-0.037	1.065	0.031	0.277	12.60
9	z[12]	-0.0	0.745	0.021	0.162	13.94
55	z[13]	0.45	0.916	0.026	0.199	13.89
56	z[14]	-0.115	0.81	0.023	0.197	7.177
57	z[15]	0.034	0.632	0.018	0.138	26.41
58	z[16]	-0.337	0.648	0.019	0.139	20.55
59	z[17]	-0.494	1.133	0.033	0.284	5.549
60	z[18]	-0.227	0.983	0.028	0.221	11.79
61	z[19]	-0.692	0.751	0.022	0.165	15.49
62	z[20]	0.027	0.725	0.021	0.161	13.48
63	z[21]	-0.027	0.808	0.023	0.146	28.26
64	z[22]	-1.134	0.825	0.024	0.203	6.681
65	z[23]	0.504	0.999	0.029	0.245	12.68
66	z[24]	0.438	0.727	0.021	0.149	19.49
67	z[25]	-0.003	0.807	0.023	0.182	9.495
68	z[26]	0.949	0.64	0.018	0.117	22.53
69	z[27]	0.402	0.863	0.025	0.2	7.688
70	z[28]	0.885	0.972	0.028	0.252	15.34
71	z[29]	0.184	0.609	0.018	0.105	27.10
72	z[30]	0.118	0.811	0.023	0.181	10.11
73	z[31]	-0.157	0.868	0.025	0.207	6.451
74	z[32]	-0.1	0.67	0.019	0.155	10.46
75	z[33]	-0.529	0.752	0.022	0.178	6.48
76	z[34]	0.437	0.63	0.018	0.114	27.99
77	z[35]	-0.067	0.657	0.019	0.137	21.14
78	z[36]	0.419	0.797	0.023	0.167	9.455
79	z[37]	-1.068	0.869	0.025	0.226	6.411
80	z[38]	0.263	0.857	0.025	0.19	14.94
81	z[39]	0.23	0.846	0.024	0.214	13.75
82	z[40]	-0.288	0.898	0.026	0.233	4.838

83	z[41]	-0.011	0.854	0.025	0.214	15.21
9						
84	z[42]	-0.088	0.739	0.021	0.157	10.40
1						
85	z[43]	-0.376	1.028	0.03	0.264	8.937
86	z[44]	-0.173	0.834	0.024	0.193	17.81
8						
87	z[45]	-0.154	0.978	0.028	0.241	7.937
88	z[46]	-0.354	0.857	0.025	0.205	20.31
4						
89	z[47]	-0.712	0.813	0.023	0.202	6.453
90	z[48]	-0.165	0.653	0.019	0.149	18.25
2						
91	z[49]	0.474	0.81	0.023	0.186	8.534
92	z[50]	0.252	0.798	0.023	0.188	8.671
93	z[51]	0.597	0.774	0.022	0.183	18.56
4						
94	z[52]	-0.438	0.952	0.027	0.206	20.02
7						
95	z[53]	-0.558	0.928	0.027	0.216	18.98
2						
96	z[54]	-0.229	0.899	0.026	0.213	10.13
2						
97	z[55]	0.112	1.193	0.034	0.312	4.819
98	z[56]	0.038	1.034	0.03	0.259	12.85
8						
99	z[57]	-0.439	0.946	0.027	0.237	11.37
9						
100	z[58]	0.47	0.93	0.027	0.246	10.75
8						
101	z[59]	0.572	1.082	0.031	0.278	7.872
102	z[60]	-0.421	0.886	0.026	0.216	15.53
3						
103	z[61]	0.056	1.036	0.03	0.267	10.55
9						
104	z[62]	0.044	0.609	0.018	0.131	7.322
105	z[63]	0.364	0.882	0.025	0.218	11.85
106	z[64]	-0.342	1.041	0.03	0.275	11.67
107	z[65]	-0.163	0.801	0.023	0.18	25.22
9						
108	z[66]	0.373	0.951	0.027	0.218	16.73
3						
109	z[67]	0.009	0.909	0.026	0.229	11.07
4						
110	z[68]	-0.487	0.798	0.023	0.198	5.812
111	z[69]	0.218	0.899	0.026	0.228	6.96
112	z[70]	0.135	0.737	0.021	0.15	19.48
2						
113	z[71]	0.748	0.869	0.025	0.222	10.89
2						
114	z[72]	-0.083	0.863	0.025	0.185	22.62
9						
115	z[73]	-0.472	0.848	0.024	0.186	10.72
7						
116	z[74]	0.095	0.853	0.025	0.212	10.35
6						
117	z[75]	-0.692	0.899	0.026	0.204	8.81
118	z[76]	-0.342	0.764	0.022	0.144	22.24
119	z[77]	0.127	1.051	0.03	0.258	7.389
120	z[78]	-0.622	0.758	0.022	0.173	12.97

2							
121	z[79]	0.624	1.112	0.032	0.289	7.533	
122	z[80]	0.312	1.161	0.034	0.295	7.763	
123	z[81]	-0.14	0.949	0.027	0.252	5.36	
124	z[82]	0.358	0.869	0.025	0.222	7.371	
125	z[83]	0.635	0.634	0.018	0.13	15.16	
4							
126	z[84]	-0.376	0.667	0.019	0.141	20.28	
127	z[85]	-0.044	0.909	0.026	0.237	14.42	
1							
128	z[86]	0.135	0.862	0.025	0.2	15.66	
129	z[87]	-0.265	0.748	0.022	0.176	13.14	
9							
130	z[88]	0.358	0.881	0.025	0.217	19.28	
1							
131	z[89]	-0.75	0.679	0.02	0.13	33.87	
132	z[90]	-0.104	0.779	0.022	0.177	16.20	
3							
133	z[91]	-0.019	0.954	0.028	0.243	7.111	
134	z[92]	0.583	0.579	0.017	0.074	42.92	
1							
135	z[93]	-0.942	0.694	0.02	0.146	15.78	
2							
136	z[94]	0.06	0.734	0.021	0.149	25.6	
137	z[95]	0.106	0.767	0.022	0.146	26.42	
8							
138	z[96]	0.637	0.798	0.023	0.196	15.69	
7							
139	z[97]	-0.348	1.003	0.029	0.259	7.034	
140	z[98]	-0.313	0.803	0.023	0.192	9.706	
141	z[99]	0.697	0.625	0.018	0.132	16.15	
142	z[100]	-0.313	0.947	0.027	0.235	6.286	

Quantiles|

Row	parameters	10.0%	50.0%	90.0%	
	Symbol	Float64	Float64	Float64	
1	article_pop_std	0.874	0.975	1.087	
2	article_score[1]	-0.084	0.316	0.687	
3	article_score[2]	0.827	1.198	1.647	
4	article_score[3]	-0.839	-0.469	-0.091	
5	article_score[4]	-1.608	-1.249	-0.89	
6	article_score[5]	-0.637	-0.276	0.06	
7	article_score[6]	-0.267	0.126	0.607	
8	article_score[7]	-1.749	-1.26	-0.931	
9	article_score[8]	0.567	0.962	1.334	
10	article_score[9]	-1.185	-0.846	-0.44	
11	article_score[10]	0.008	0.451	0.759	
12	diffcp[1]	0.085	0.101	0.118	
13	diffcp[2]	0.084	0.1	0.118	
14	diffcp[3]	0.084	0.101	0.117	
15	diffcp[4]	0.09	0.106	0.122	
16	diffcp[5]	0.083	0.097	0.116	
17	diffcp[6]	0.081	0.095	0.111	
18	diffcp[7]	0.082	0.098	0.116	
19	diffcp[8]	0.084	0.102	0.118	
20	diffcp[9]	0.084	0.101	0.117	
21	diffcp[10]	0.081	0.097	0.117	
22	reviewer_gain[1]	0.974	1.0	1.03	
23	reviewer_gain[2]	0.968	0.999	1.03	
24	reviewer_gain[3]	0.969	0.999	1.03	

25	reviewer_gain[4]	0.973	1.0	1.027	
26	reviewer_gain[5]	0.97	0.999	1.029	
27	reviewer_gain[6]	0.969	0.999	1.026	
28	reviewer_gain[7]	0.975	1.001	1.029	
29	reviewer_gain[8]	0.969	0.998	1.028	
30	reviewer_gain[9]	0.971	1.002	1.027	
31	reviewer_gain[10]	0.971	0.998	1.03	
32	reviewer_noise[1]	0.525	0.769	1.059	
33	reviewer_noise[2]	0.516	0.784	1.135	
34	reviewer_noise[3]	0.53	0.761	1.084	
35	reviewer_noise[4]	0.507	0.741	1.098	
36	reviewer_noise[5]	0.521	0.735	1.072	
37	reviewer_noise[6]	0.525	0.756	1.076	
38	reviewer_noise[7]	0.546	0.784	1.118	
39	reviewer_noise[8]	0.527	0.755	1.09	
40	reviewer_noise[9]	0.539	0.78	1.082	
41	reviewer_noise[10]	0.524	0.767	1.084	
42	$r\sigma$	0.544	0.77	1.065	
43	z[1]	-1.151	-0.31	0.795	
44	z[2]	-0.831	0.723	1.872	
45	z[3]	-0.949	-0.156	1.014	
46	z[4]	-1.275	-0.223	1.704	
47	z[5]	-1.759	-0.784	0.551	
48	z[6]	-1.35	-0.632	1.286	
49	z[7]	-0.465	0.323	1.083	
50	z[8]	-0.494	0.563	1.585	
51	z[9]	-0.302	0.487	1.516	
52	z[10]	-0.534	0.379	1.76	
53	z[11]	-1.173	-0.147	1.569	
54	z[12]	-1.014	0.07	0.842	
55	z[13]	-0.751	0.38	1.796	
56	z[14]	-1.243	-0.123	1.012	
57	z[15]	-0.771	0.031	0.853	
58	z[16]	-1.25	-0.297	0.442	
59	z[17]	-2.104	-0.306	0.875	
60	z[18]	-1.554	-0.198	0.966	
61	z[19]	-1.665	-0.61	0.241	
62	z[20]	-1.008	0.041	0.878	
63	z[21]	-1.14	0.061	0.878	
64	z[22]	-2.298	-1.079	-0.123	
65	z[23]	-0.879	0.522	1.74	
66	z[24]	-0.549	0.471	1.395	
67	z[25]	-1.006	-0.051	1.104	
68	z[26]	0.146	0.922	1.803	
69	z[27]	-0.679	0.357	1.429	
70	z[28]	-0.298	0.8	2.007	
71	z[29]	-0.61	0.201	1.0	
72	z[30]	-0.827	0.07	1.237	
73	z[31]	-1.334	-0.06	0.918	
74	z[32]	-0.96	-0.144	0.814	
75	z[33]	-1.393	-0.623	0.675	
76	z[34]	-0.457	0.477	1.266	
77	z[35]	-0.873	-0.016	0.738	
78	z[36]	-0.504	0.366	1.5	
79	z[37]	-2.279	-0.977	-0.048	
80	z[38]	-0.953	0.315	1.371	
81	z[39]	-0.917	0.28	1.27	
82	z[40]	-1.37	-0.373	1.006	
83	z[41]	-1.079	0.08	0.992	
84	z[42]	-0.975	-0.086	0.865	
85	z[43]	-1.597	-0.542	1.034	
86	z[44]	-1.293	-0.106	0.858	
87	z[45]	-1.436	-0.343	1.214	
88	z[46]	-1.511	-0.422	0.794	
89	z[47]	-1.73	-0.769	0.444	
90	z[48]	-1.111	-0.149	0.736	
91	z[49]	-0.664	0.603	1.46	
92	z[50]	-0.838	0.281	1.317	
93	z[51]	-0.452	0.73	1.523	
94	z[52]	-1.613	-0.449	0.867	
95	z[53]	-1.427	-0.556	0.547	
96	z[54]	-1.312	-0.37	0.98	
97	z[55]	-1.687	0.249	1.659	
98	z[56]	-0.981	-0.174	1.628	
99	z[57]	-1.525	-0.439	0.791	

100	z[58]	-0.772	0.454	1.715	
101	z[59]	-0.685	0.422	2.226	
102	z[60]	-1.804	-0.218	0.519	
103	z[61]	-1.508	0.192	1.232	
104	z[62]	-0.8	0.112	0.78	
105	z[63]	-0.944	0.454	1.501	
106	z[64]	-1.488	-0.413	0.666	
107	z[65]	-1.153	-0.221	0.916	
108	z[66]	-1.045	0.334	1.528	
109	z[67]	-1.084	0.011	1.131	
110	z[68]	-1.452	-0.517	0.542	
111	z[69]	-0.935	0.213	1.391	
112	z[70]	-0.788	0.065	1.179	
113	z[71]	-0.425	0.751	1.935	
114	z[72]	-1.181	-0.146	1.056	
115	z[73]	-1.491	-0.574	0.815	
116	z[74]	-0.961	0.019	1.229	
117	z[75]	-1.873	-0.738	0.501	
118	z[76]	-1.326	-0.435	0.62	
119	z[77]	-0.942	-0.03	1.759	
120	z[78]	-1.603	-0.685	0.453	
121	z[79]	-0.747	0.486	2.174	
122	z[80]	-1.375	0.647	1.632	
123	z[81]	-1.283	-0.27	1.155	
124	z[82]	-0.685	0.19	1.668	
125	z[83]	-0.182	0.594	1.434	
126	z[84]	-1.35	-0.309	0.433	
127	z[85]	-1.218	-0.085	1.201	
128	z[86]	-0.992	0.041	1.265	
129	z[87]	-1.23	-0.18	0.524	
130	z[88]	-0.731	0.291	1.432	
131	z[89]	-1.607	-0.781	0.171	
132	z[90]	-1.068	-0.183	1.009	
133	z[91]	-1.423	0.045	1.185	
134	z[92]	-0.148	0.598	1.296	
135	z[93]	-1.827	-0.933	-0.054	
136	z[94]	-1.029	0.158	0.987	
137	z[95]	-0.872	0.091	1.142	
138	z[96]	-0.351	0.638	1.652	
139	z[97]	-1.49	-0.398	0.953	
140	z[98]	-1.329	-0.268	0.686	
141	z[99]	-0.179	0.731	1.426	
142	z[100]	-1.584	-0.251	0.917	

dc[1].df.r_hat

142-element Array{Any,1}:

```

1.004
1.049
1.04
1.07
1.0
1.001
1.182
1.009
1.001
1.089
⋮
1.023
1.08
1.045
1.002
1.008
1.001
1.071
1.101
1.403

```

Maximum r_hat

`maximum(filter(isfinite, dc[1].df.r_hat))`

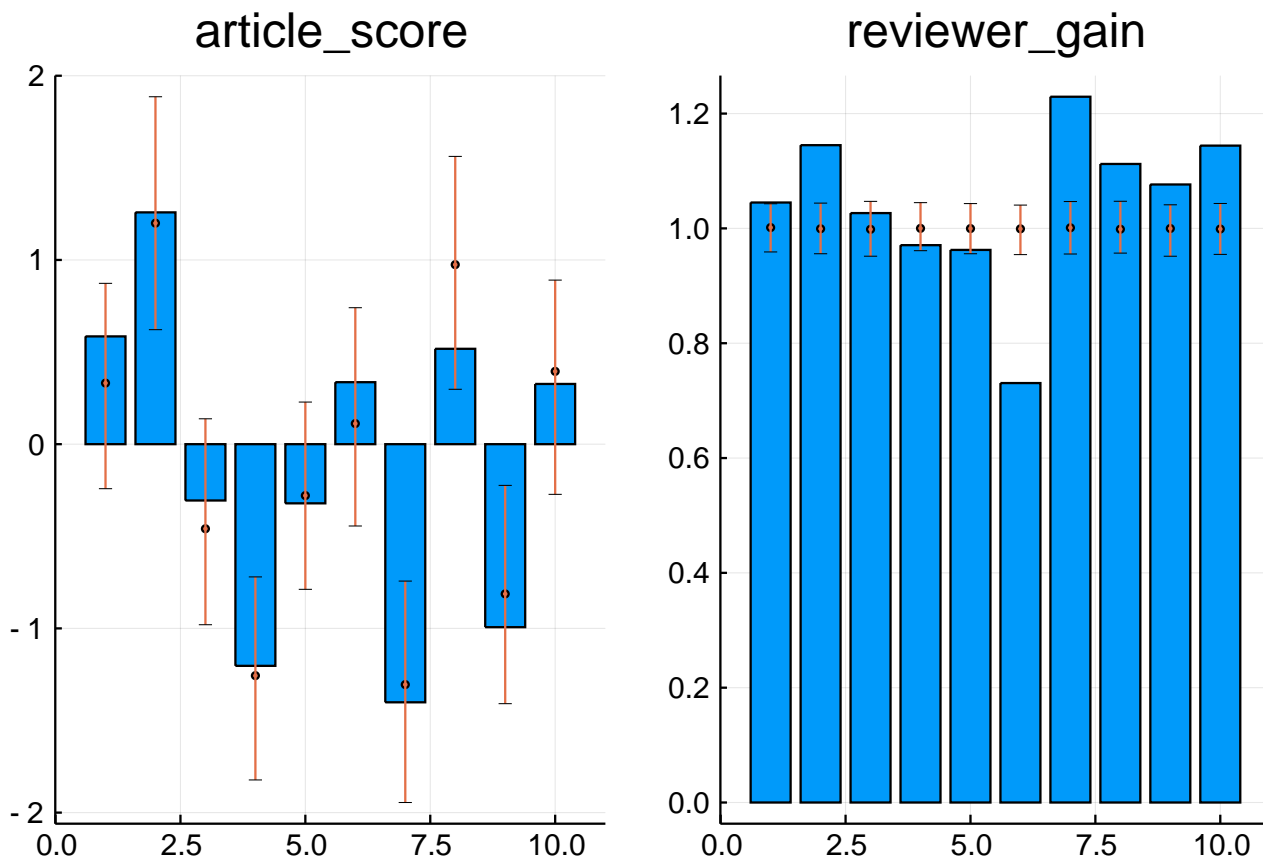
1.585

`median(filter(isfinite, dc[1].df.r_hat))`

1.0365

We plot the same figure of the posterior article scores and reviewer gains as we did for Soss

```
figs = map((:article_score,:reviewer_gain)) do s
  bar(getproperty(truth, s))
  prop = getproperty(p, s)
  errorbarplot!(1:length(prop), prop, seriestype=:scatter, legend=false, title=string(s), m=2)
end
plot(figs...)
```




When sampling with turing, the posterior for reviewer gain is very different from when sampling with Soss

We also check how the model fared when estimating the rank of the articles

```
lo = logodds(truth.Rv)
observed_score = map(1:na) do j
  median([lo[ind] for ind in eachindex(indv) if indv[ind][2] == j])
end
print_rank_results(truth, p, observed_score)
```

Percentage of correct rank from model 0.4
 Percentage of correct rank from observed score 0.2
 Rankdist between correct rank and model 0.8
 Rankdist between correct rank and observed score 1.0

Article scores



3 MAP

using Turing and Optim, we perform MAP estimation to try to figure out why the inference above is shaky

```
function get_nlogp(model)
  vi = Turing.VarInfo(model)
  function nlogp(sm)
    spl = Turing.SampleFromPrior()
    new_vi = Turing.VarInfo(vi, spl, sm)
    try # If the call below fails, we just return a large value
      model(new_vi, spl)
    catch
      return 1e4
    end
    -new_vi.logp
  end
  return nlogp
end
```

get_nlogp (generic function with 1 method)

Define our data points.

```
model = cum_model(indv, truth.Rv)
nlogp = get_nlogp(model)
```

nlogp (generic function with 1 method)

We set the initial parameter estimate the true generating parameters

```
using Optim
p0 = [truth.r0;
      truth.article_pop_std;
      truth.reviewer_noise;
      truth.reviewer_gain;
      truth.article_score;
```

```
truth.diffcp;
truth.z]
```

```
nlogp(p0)
```

```
result = Optim.optimize(nlogp, p0, GradientDescent(), Optim.Options(store_trace=true, show_trace=false, show_every=1, iterations=3000, allow_f_increases=false,
time_limit=300, x_tol=0, f_tol=0, g_tol=1e-8, f_calls_limit=0, g_calls_limit=0), autodiff=:forward)
```

```
* Status: success
```

```
* Candidate solution
```

```
Minimizer: [1.69e-13, 9.45e-01, 6.53e-03, ...]
```

```
Minimum: 4.258357e+02
```

```
* Found with
```

```
Algorithm: Gradient Descent
```

```
Initial Point: [1.02e-03, 9.45e-01, 6.53e-03, ...]
```

```
* Convergence measures
```

```
|x - x'| = 0.00e+00 ≤ 0.0e+00
```

```
|x - x'|/|x'| = 0.00e+00 ≤ 0.0e+00
```

```
|f(x) - f(x')| = 0.00e+00 ≤ 0.0e+00
```

```
|f(x) - f(x')|/|f(x')| = 0.00e+00 ≤ 0.0e+00
```

```
|g(x)| = 4.73e+12 ≰ 1.0e-08
```

```
* Work counters
```

```
Seconds run: 0 (vs limit 300)
```

```
Iterations: 2
```

```
f(x) calls: 112
```

```
f(x) calls: 112
```

Using ForwardDiff we can have a look at the gradient and Hessians at the solution to the optimization problem

```
using ForwardDiff
```

```
@time H = ForwardDiff.hessian(nlogp, result.minimizer)
```

```
30.716750 seconds (25.64 M allocations: 1.247 GiB, 3.68% gc time)
```

```
142×142 Array{Float64,2}:
```

```
-2.79287e25 -0.0 -100.0 -100.0 ... -0.0 -0.0 -0.
0
-0.0 175.862 -0.0 -0.0 -0.0 -0.0 -0.
0
-100.0 -0.0 103.593 -0.0 -0.0 -0.0 -0.
0
-100.0 -0.0 -0.0 104.882 -0.0 -0.0 -0.
0
-100.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-100.0 -0.0 -0.0 -0.0 ... -0.0 -0.0 -0.
0
-100.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-100.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-100.0 -0.0 -0.0 -0.0 -0.073255 -0.0 -0.
0
⋮ ⋮ ⋮
-0.0 -0.0 -0.0 -0.625056 -0.0 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 ... -0.0 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 -0.0 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 1.00355 -0.0 -0.
0
-0.0 -0.0 -0.0 -0.0 ... -0.0 1.00084 -0.
0
```

```
-0.0 -0.0 -0.0 -0.0 -0.0 -0.0 1.  
0
```

```
g = ForwardDiff.gradient(nlogp, result.minimizer)
```

```
142-element Array{Float64,1}:  
4.726837693210614e12  
-2.8891193081493203  
4.0986373087057935  
-0.3228856437743466  
-1.1804530149784491  
20.90811582135648  
0.6423516061180331  
7.897144923755578  
-0.9050132595675503  
11.03459806864493  
⋮  
2.3702408023354438  
-0.8017933986572475  
1.5385458497150675  
0.3523061278826322  
-0.014117655279035163  
-2.3884446519452167  
1.4191179001315486  
1.61543658807263  
-1.5205557236621878
```

```
cond(H)
```

```
2.923568656311296e25
```

This page was generated using [Literat.jl](#).